

Jamie Park

jamie.park@email.com | (734) 555-0192 | linkedin.com/in/jamiepark | github.com/jamiepark

EDUCATION

University of Michigan - B.S. Computer Science & Statistics

Aug 2020 - May 2024

GPA: 3.72 / 4.00

Coursework: Machine Learning, Deep Learning, Statistical Inference, Database Systems, Data Structures & Algorithms, Probability Theory

WORK EXPERIENCE

Data Analytics Intern - ShipFast Logistics

May 2023 - Aug 2023

- Built 6 Tableau dashboards tracking delivery KPIs across 12 regional hubs, reducing manual reporting time by 40%
- Wrote SQL queries against a PostgreSQL database to extract and aggregate shipment delay data for weekly operations review
- Automated weekly Excel summary reports using Python and openpyxl, eliminating a 3-hour recurring manual process

PROJECTS

Sentiment Classifier (Undergraduate Thesis)

Jan 2024 - Apr 2024

- Fine-tuned a BERT model on 50,000 Amazon product reviews for multi-class sentiment classification (Positive/Neutral/Negative)
- Achieved 91.3% accuracy on held-out test set, outperforming TF-IDF + Logistic Regression baseline (82.1%)
- Stack: Python, PyTorch, HuggingFace Transformers, Scikit-learn, Jupyter Notebook

Movie Recommendation System

Sep 2023 - Dec 2023

- Implemented collaborative filtering via matrix factorization (ALS) trained on the MovieLens 1M dataset
- Exposed recommendations through a REST API; evaluated with RMSE and precision@k metrics
- Stack: Python, NumPy, Pandas, Scikit-learn, Flask, Git

Campus Event Multi-Label Classifier

Feb 2023 - Apr 2023

- Built a multi-label text classifier to auto-tag university event listings using Logistic Regression and Random Forest
- Scraped and cleaned 4,000 event records; achieved F1-score of 0.84 across 8 categories
- Stack: Python, Scikit-learn, BeautifulSoup, Pandas

SKILLS

Languages:	Python, R, SQL, Bash
ML / AI:	PyTorch, HuggingFace Transformers, Scikit-learn, NumPy, Pandas, Matplotlib
Data & Viz:	Tableau, Jupyter Notebook, PostgreSQL, openpyxl
Tools:	Git, GitHub, VS Code, Linux