Jamie Park

jamiepark359@gmail.com | (734) 555-0192 | linkedin.com/in/jamiepark | github.com/jamiepark

EDUCATION

University of Michigan - B.S. Computer Science & Statistics Aug 2020 - May 2024

GPA: 3.72 / 4.00

Coursework: Machine Learning, Deep Learning, Statistical Inference, Database Systems, Data Structures & Algorithms, Probability Theory, Object-Oriented Programming (Java)

WORK EXPERIENCE

ML Engineering Intern - ShipFast Logistics May 2023 - Aug 2023

- Trained and deployed a shipment delay prediction model (XGBoost) to AWS SageMaker; model served live predictions consumed by the operations dashboard
- Built a LangChain-powered internal Q&A tool over logistics docs using OpenAI embeddings and a Pinecone vector store; reduced analyst lookup time by ~60%
- Wrote SQL queries against a PostgreSQL database to build training features from 2M+ shipment records; automated weekly reporting with Python and openpyxl
- Stack: Python, XGBoost, Scikit-learn, LangChain, Pinecone, AWS SageMaker, PostgreSQL

PROJECTS

Sentiment Classifier (Undergraduate Thesis) Jan 2024 - Jun 2024

- Fine-tuned a BERT model on 50,000 Amazon product reviews for multi-class sentiment classification (Positive/Neutral/Negative)
- Achieved 91.3% accuracy on held-out test set, outperforming TF-IDF + Logistic Regression baseline (82.1%)
- Tracked experiments with MLflow Model Registry; logged training curves and W&B sweeps for hyperparameter tuning
- Containerised with Docker and deployed to GCP Cloud Run; GitHub Actions CI/CD pipeline redeploys on every push to main
- Stack: Python, PyTorch, HuggingFace Transformers, MLflow, Weights & Biases, Docker, GCP Cloud Run, GitHub Actions

RAG Document Assistant - LlamaIndex + Azure Jan 2025 - Mar 2025

- Built a retrieval-augmented generation (RAG) system over 500+ technical PDFs using LlamaIndex with a ChromaDB vector store
- Deployed the FastAPI backend and ChromaDB instance to Azure Container Apps; used Azure Blob Storage for document ingestion
- Evaluated retrieval quality with RAGAS (faithfulness, answer relevancy); iterated on chunking strategy to improve faithfulness score from 0.71 to 0.89
- Stack: Python, LlamaIndex, LangChain, ChromaDB, Azure Container Apps, Azure Blob Storage, FastAPI, RAGAS

End-to-End MLOps Pipeline - Kubeflow + MLflow Sep 2024 - Dec 2024

- Built a retrainable ML pipeline using Kubeflow Pipelines: data validation, feature engineering, training, and evaluation as separate components
- Integrated MLflow Model Registry as the promotion gate; automated retraining triggered by data-drift detection (Evidently)
- Deployed via blue/green rollout on Kubernetes; wrote a Java Spring Boot microservice for batch prediction serving
- Stack: Python, Kubeflow Pipelines, MLflow, Evidently, Docker, Kubernetes, Java, Spring Boot

SKILLS

Languages:	Python, R, SQL, Bash, Java
ML / AI:	PyTorch, HuggingFace Transformers, Scikit-learn, XGBoost, NumPy, Pandas
LLM / RAG:	LangChain, LlamaIndex, OpenAI API, ChromaDB, Pinecone
MLOps:	MLflow, Weights & Biases, Kubeflow Pipelines, Evidently, SageMaker

Cloud:	AWS (SageMaker, S3), GCP (Cloud Run), Azure (Container Apps, Blob Storage)
Infra:	Docker, Kubernetes, GitHub Actions, Spring Boot
Data & Viz:	Tableau, Jupyter Notebook, PostgreSQL, openpyxl